


16 experts were sure AI sped them up — and got the sign wrong



Against Expert Forecasts and Developer Self-Reports, Early-2025 AI Slows Down Experienced Open-Source Developers



In this RCT, 16 developers with moderate AI experience complete 246 tasks in large and complex projects on which they have an average of 5 years of prior experience.



Predicted -24% · believed -20% · actual +19% time

METR experiment, first half of 2025

16 experienced open-source developers · 246 real tasks in their own familiar repositories; measured real time, not perception

Three numbers about one and the same thing

Predicted before the experiment: AI will speed up by -24%

Believed after working: sped up by roughly -20%

Measured actual: with AI, tasks took +19% longer

— METR RCT, n=16 experienced OSS developers, 246 tasks, 2025

The professionals got the magnitude right but the sign wrong. "It feels like the tool helps" is a hypothesis, not data; reliability comes from discipline with built-in verification, not from perception. ^[1]

[1] METR — RCT: AI and experienced OSS developers (+19% time)

04

LECTURE 4

AI across the software development lifecycle

Course "Deliberate use of AI" · software developers



Reliability comes not from the tool but from engineering discipline by phase: which artifact and which check each one needs.



We stand on the foundation of Module 1 — take one industry and break it down by phase



From Lecture 1

The layered picture "model → chat → agent → application"; a prompt as "role + task + context". We use it as a vocabulary.



From Lecture 2

Why AI produces "almost right" text: the answer is generated token by token, so plausible ≠ correct. In one phrase.



From Lecture 3

The complexity ladder ("stay on the lowest rung") + the criterion "when not AI at all" — a deterministic task is solved by ordinary code.



From Lecture 3

The agent loop plan → act → check → iterate (4 points of failure) + prompt injection (defense is architectural, not by filtering).

We carry this over intact and don't re-explain it — we only apply it to software development. Module 1 gave the apparatus for choosing an architecture; Lecture 4 shows which discipline makes AI in one industry reliable by phase — and where it breaks without it.

The lecture's question is about discipline, not about the tool

AI writes code ever better — but what makes AI development reliable?

Not the tool, but engineering discipline by phase: which human-owned artifact and which check each phase needs, where the methods of different players converge — and what is not delegated to the tool in any phase. ^[1]

Which tool is best?

Goes stale in a quarter; "best" can't be named without "for what, in which phase, in which mode".



Which practice is warranted in the phase?

Which artifact does it end in? Where is a human mandatory? — stable for years.

The tool is secondary, the practice primary. When answering, an engineer names not a logo but the phase's appropriate practice, its artifact, and the point of mandatory human control — knowledge that will outlive the swap of any vendor. ^[2]

This lecture is a digest of practices: modern approaches of the leaders + time-tested classics



A. Modern practices (leaders)

[1] [Anthropic — Claude Code / AI-Native SDLC](#)

agentic git loop: each stage commits an artifact; human at the gates

[2] [OpenAI — Model Spec](#)

spec-as-contract; clause = example prompt = test

[3] [GitHub — Spec Kit](#)

"intent is the source of truth"; small verifiable tasks

[4] [Google — DORA 2025](#)

"AI amplifies what is already there"; seven capabilities

[5] [Thoughtworks — Exploring Gen AI](#)

the assistant suggests, the developer owns; harness engineering

[6] [Willison — Vibe engineering](#)

the disciplines an LLM rewards: tests, plans, reviews



B. Time-tested classics

[7] [Brooks — No Silver Bullet](#)

essential vs accidental complexity; "what to build" is the human's

[8] [Beck — TDD: By Example](#)

red-green-refactor; test-as-specification

[9] [Nygard — ADR](#)

immutable "why" records under version control

[10] [Ford/Parsons — Evolutionary Architectures](#)

fitness functions: "fitness" is objective and automatic

[11] [Fowler — Refactoring](#)


the discipline of small verifiable changes

[12] [Brown — C4](#)

architecture-as-code: a textual, diffable model for AI

**Tools will change over quarters — these practices are stable: they rest on the nature of complexity, not on product maturity.
All names are clickable links to the primary sources.**

Discipline is a cycle of phases, and each phase is owned by a human through its artifact



A phase is a stage with its own input / output / artifact.

AI works INSIDE the nodes (drafts requirements, proposes an ADR, writes code in a PR), but each node is owned by a human — who reads, edits, accepts, is accountable ^[1].

The practices mature unevenly

strong: requirements, implementation, testing, review · thin: architecture, delivery, operations · bright spot: documentation.

Anthropic ^[1], OpenAI ^[2], DORA ^[3], and Thoughtworks ^[4] independently converged on this phase skeleton — so what we have is a method, not a fashion.

Autonomy is a property of the mode, not the brand

The autonomy ladder — a supporting lens



- A autocomplete**
finishes the line; the human accepts each one
- B small tasks**
a function/fix in a dialog; the human reviews after
- C coding agent**
plans, edits files, runs tests; the human does the merge
- D orchestrator**
takes a task from the tracker → PR; the human is the prod gate

Mode $\neq$ brand — one product lives on several rungs at once ^[2]

Copilot — A, B, C and D · Cursor — Tab (A), Cmd-K (B), Composer (C) · Claude Code — C, rising to D

B ↔ C does it iterate and run tests without you? (yes → C)
C ↔ D where the task comes from, where it lands (tracker → PR → D)

The higher the autonomy, the stricter the criterion "a human is mandatory here" and the tighter the harness. "We use Copilot" reports neither the level nor the phase — name the mode and the phase, not the logo. ^[1]

The method decides, the tool executes

A methodological practice is the decision about which artifact to produce, in what order, with what check, and who is accountable. The tool is the executor: interchangeable and secondary.

Volatility

Tools and their "leadership" change over quarters; practices are stable for years — they rest on the nature of complexity (Brooks ^[1]), not on product maturity.

The DORA multiplier

"AI amplifies what is already there" ^[2]: build the discipline first, not pick the tool — a tool without discipline multiplies chaos.

Accountability

The tool is not accountable for consequences, the human is; accountability materializes in artifacts and gates, that is, in the practice.

"We adopted an AI tool" $\neq$ "we adopted an AI discipline". The tool is there, the practice is not — and the DORA multiplier works the wrong way. Behind this stand all the failures of the lecture: prompt-and-pray, poisoned context, the 70% problem, the Replit incident — everywhere the tool applied without the practice.

Method-first order: first the phase's practice (artifact, gate, who is accountable), then the tool to fit it. The tool is chosen to fit the discipline, not the discipline to fit the tool. ^[3]

SECTION 1

Requirements — the first artifact, before any code

The cycle's first phase turns a vague intent into requirements. What leads here is not the tool but the discipline of requirements (spec-driven): between intent and code a human-reviewable artifact is placed, and the decision "what exactly we need" stays human.



0

Intro

1

Requirements

2

Architecture

3

Implementation

4

Testing

5

Review+Sec.

6

Delivery+

7

Synthesis

The first development artifact is requirements, not code



Repository: requirements are versioned next to the code ^[3]



requirements.md (spec.md)

what the system must do



design.md

constraints and decisions



tasks.md

decomposition into tasks



src/...

code — generated from requirements

The order is enforced: requirements → design → tasks (Kiro / Spec-Kit ^[3]). The costliest errors are in this phase.

OpenAI Model Spec ^[2]

living, versioned Markdown: requirements are a durable, reviewable, diffable artifact, not fleeting prompts.

Sean Grove, OpenAI · "The New Code" ^[5]

"the source requirements — that is the valuable artifact"; code is "structured communication".

Martin Fowler ^[4]

the bottleneck of AI development is intent: the model writes code ever better, but expressing what to build is still hard.

Grove's "code ~10-20% of the value" is a rhetorical provocation, not an empirical measurement: its job is to shift the focus to requirements.

AI is strong at structuring and completing intent; the intent itself — what the system needs — is decided by the human ^[1].

How to run requirements: structure (how to write) + process (how to maintain)

✓≡ STRUCTURE — how to write requirements

- **Stories + acceptance criteria** ^[3]

"As a <role>, I want <goal>, so that <benefit>" + verifiable criteria for each story.

- **EARS notation** ^[7]

"WHEN <trigger>, the system SHALL <response>" (Mavin 2009) — removes "should/may", makes a requirement testable.

- **Functional vs non-functional** ^[6]

behavior separate from characteristics (latency / cost / security); NFRs are enforced by fitness functions.

- **requirements → design → tasks** ^[3]

the enforced order of 3 files (Kiro / Spec-Kit); DoD — small independently testable units.

🔄 PROCESS — how to maintain requirements

- **Elicitation: the interrogatory LLM** ^[4]

the model ASKS questions (Fowler's "Interrogatory LLM"), surfacing unstated assumptions — instead of "prompt-and-pray".

- **Review and sign-off BEFORE code** ^[1]

requirements are reviewed and signed off by a human before generation; accept/reject = "the merge".

- **Versioning next to the code** ^[3]

requirements are diffable Markdown in the repository, not in a wiki / chat; a durable artifact, not a fleeting prompt.

- **Syncing with change** ^[9]

keep current like an ADR; the human owns "what to build", AI owns structure and completeness.

Durable pattern: EARS + decomposition + requirements-as-check + human sign-off — will outlive any tool. Hype: "our pipeline of commands = requirements discipline". Tools execute, the method leads.

prompt-and-pray: the bug is not in the code, but in a requirement nobody checked



The visible 'works in the demo' is the tip; below the water are dozens of unstated assumptions the model defaulted on.

prompt-and-pray — one vague prompt ('build me a booking system') and hope. This skips the discipline: no requirements artifact, no human checkpoint between intent and code.

The model silently fills in the decisions: a booking in the past? overlapping bookings? who cancels someone else's? time zones? — for each it takes a plausible default. It 'works' in the demo, breaks on the first real conflict.

The code is correct relative to what the model assumed. The bug is not in the code — it is that nobody checked the assumptions; they are invisible in the code.

The mirror-image extreme (Encarnacao, 'The Emperor's New Code'): 'the spec = the single truth, you need not read the code'. But a spec underspecifies behavior; 'I'll regenerate from the spec' is a new guess, not the same product. The code remains the source of truth.

The bottleneck is not the model's ability to write code, but the precision of stating intent (essential complexity, Brooks ^[1]). The alternative is not 'no AI', but restoring the human checkpoint: requirements accepted before code ^[2].

SECTION 2

Architecture — before code, and it must be managed

After requirements comes not code straight away, but architecture: deciding what to assemble the system from. This is essential complexity, led by the human; the leading practices — ADR, fitness functions, architecture-as-code — teach you to manage it with AI, not delegate it to AI.



0

Intro

1

Requirements

2

Architecture

3

Implementation

4

Testing

5

Review+Sec.

6

Delivery+

7

Synthesis

After requirements comes architecture, not code straight away

what is needed

requirements



what to assemble from

architecture

**cannot
skip**



how to write

code

The phase's output is a small number of hard, hard-to-reverse forks: component boundaries, the data model, the priority of quality attributes.

Jump straight to code →

Architecture erosion — the gap between what was intended and what was implemented, a decay in maintainability.

Codebase cognitive debt (Thoughtworks Radar, Hold ring ^[2]): the gap between how the system is built and the team's understanding — it 'lives in people's heads', not in artifacts. The remedy Radar names — architectural fitness functions ^[3].

'Deciding what to build' is essential complexity (Brooks, 'No Silver Bullet', 1986 ^[1]): a trade-off choice is not delegated. AI is useful only at the periphery — options, explaining a pattern, a draft diagram.

Four practices for managing architecture with AI (tools are secondary)

 ADR	 Fitness function	 C4 / arch-as-code	 Evolutionary arch.
WHAT IT IS Half a page of an immutable record per decision (context·decision·status·consequences); stores the 'why'	Auto-check of an architectural attribute on every commit ('payment does not depend on the UI'; 'response < 200 ms')	Architecture machine-readably (C4: Context/Container/Component/Code; DSL — PlantUML/Mermaid/Structurizr)	ADR + fitness functions + arch-as-code together = incrementality + managed change
WHO PRESCRIBES IT Nygard 2011 ^[1] ; Radar — ADOPT ^[4]	Thoughtworks; Rebecca Parsons ^[4]	Simon Brown (C4) ^[3] ; Structurizr	Ford, Parsons, Kua ^[2]
AI ROLE (SECONDARY) edits, cross-checks — but the human decides and justifies	convenient for writing fitness functions; they also validate generated code	reads as context, generates diagrams; drift detection — Structurizr (model vs code)	executes inside each of the three practices
WHERE THE HUMAN IS the fork's author = the ADR's author	decides which invariant is critical	owns the textual model	holds the direction of evolution

The durable pattern: automatic architectural control on every commit. Vendor hype: 'our product will ensure the architecture itself'. The human owns the 'why', AI codes and checks.

Poisoned context: AI does not tell 'how it ended up' from 'how it's right'

This happens **WHEN** the architecture is not described and there is no process for managing it. With practices in place (ADR, fitness functions, arch-as-code) the loop breaks.

The poisoning loop (Böckeler 2026, Thoughtworks) ^[1]

● bad design

AI copies ('how it's done here')

worse design

AI copies even more confidently

Böckeler, honestly: 'we don't yet have a good way to mitigate this'.



The human owns the forks

makes the architectural decisions; AI at the periphery under human choice.



ADR ^[2]

human-written context 'we decided X because Y, rejected Z' — shared understanding against poisoning.



Fitness functions + modular code ^[3]

deterministic invariants break the loop; clear components give managed context.

AI sees a pattern and continues it — it does not tell a good example from a bad one. The worse the existing architecture, the more strongly AI entrenches it.

SECTION 3

Implementation — discipline and harness

Here AI writes code, and the phase is strong — but strong under discipline. Three practices hold reliability: split into small verifiable units, maintain a persistent memory layer in the repository, and surround the model with a deterministic harness.



0

Intro

1

Requirements

2

Architecture

3

Implementation

4

Testing

5

Review+Sec.


6

Delivery+

7

Synthesis

Working discipline: small verifiable units + the explore→plan→code→commit loop



The order is enforced: generation before exploration and a plan is prompt-and-pray at the code level. — the explore→plan→code→commit loop, Anthropic ^[1]

Small verifiable units


Each piece is implemented and verified in isolation: AI gets a deterministic self-check, the human gets a small diff that can actually be reviewed. Osmani ^[2]: the smaller the AI's proposal, the more real the review; a giant diff goes unread.

AI takes the accidental complexity (boilerplate, a routine handler). **The human takes the essential**: what we're building, what's risky, what's correct, whether it can merge. ^[3]


AI takes part in two philosophies — in the editor (synchronously) and asynchronously (in isolation → PR); this is a property of the mode, we'll cover it secondarily.

The discipline of the loop and the small diff is not bureaucracy, but a way to keep AI in the zone where the human really controls the result.

A persistent memory layer in the repository — what the agent reads every session



A deterministic scaffold-gate around a non-deterministic model



The 70% problem: AI speeds up the first 70%, but not the last 30% — understanding

first ~70% — fast, cheap

last 20-30%

The last 20-30% — edge cases, error handling, security, integration, behavior under load — stay just as hard and require senior oversight. The gap is structural: the system's specifics are absent from the training data.

'Almost right' code is costlier than obviously wrong code: it passes a quick glance and breaks on an edge case. Work shifts from writing to debugging someone else's plausible logic.

Stack Overflow 2025: 66%

of developers named their top frustration 'solutions that are almost right, but not quite'.

GitClear · 211M lines, 2020-2024 ^[2]

clones 8.3% → 12.3%; refactored ~25% → <10%; churn 3.3% → 5.7%. (Correlation, not an RCT.)

The knowledge paradox (Osmani) ^[1]

seniors challenge the AI's output, juniors accept it ('a house of cards') — AI amplifies the experienced more.

The alternative — small verifiable units + a harness + reading the diff before accept; duplication and churn metrics in CI as a gate. Merge is always the human.

A brand and a benchmark number != engineering discipline



Verified (~500 tasks, public code) — top ~88-89%. Pro (private, contamination-resistant) — leader ~64%. Gap ~24 pp: trust in the number is inversely proportional to the unfamiliarity and criticality of your task.

Devin (Cognition): 13.86% ^[1]

vs a 1.96% baseline — but only on 25% of the bench (79 of 570 tasks), acknowledged contamination, 45-min limit; independently ~15% (3 of 20).

OpenAI: "~80% Verified" / "70% more PRs"

OpenAI itself: ~59% of "failures" are test-design defects, not the model's; "70% more PRs" — with no denominator.

Cursor: Composer "frontier, 4x faster"

its own blog admits: GPT-5 and Sonnet 4.5 "both outperform" → frontier-fast, not frontier-best.

Five questions for any vendor number:

1. Which slice?
2. Contamination?
3. Comparison baseline?
4. Fact or marketing?
5. What's in the fine print? ^[2]

Devin 13.86% — technically true on exactly a quarter of the tasks. A number can be true and misleading; a high figure doesn't answer the merge-gate question. A brand/benchmark doesn't replace discipline.

SECTION 4

Testing — TDD as a discipline

Implementation produces code — testing produces a verified claim about its correctness. TDD discipline leads here: the test is an executable specification, subject to neither the "almost right" nor the perception gap.



0

Intro

1

Requirements

2

Architecture

3

Implementation

4

Testing

5

Review+Sec.

6

Delivery+

7

Synthesis

TDD-as-approach: the human decides what to check; the run is deterministic

The red-green-refactor cycle (Kent Beck, TDD) ^[1] — the human owns the test spec

 **red** a failing test expresses a requirement

green the code that makes it pass

refactor improve while keeping it green

↑ *repeats*

AI writes tests fast — volume (accidental). **The human decides WHAT the test must assert** — essential.

Verification is not outsourced to the model

Willison / Fowler: "if you haven't seen it work, it's not a working system." Tests are run by a deterministic executor (script / CI), not by the model's word. Incident → permanent regression test.

An important nuance — structure != ritual

The value of TDD is the structure (spec-test + gate), not the ritual of forcing the order on the agent. Böckeler ^[2]: TDD-first in the agent loop — no gain + ~3x tokens ("I stopped telling agents to write tests first").

Tests-as-guardrails (Fowler) ^[3]: a test forces the interface without coupling to the implementation — which is why the TDD structure is valuable.

Executors (secondary): AWS Q /test · Qodo · JetBrains Junie · Anthropic (failing test → fix + Stop-hook as a gate).

Durable pattern: test-as-executable-specification + a deterministic run gate. Hype: "AI covered the code with tests on its own."

Green tests and high coverage can lie — the gate must be honest




"all green" lies (Fowler) ^[1]

"An LLM will happily say 'all tests green' even when there are failures." The mechanism is the same — the model generates a plausible report with the same token-by-token sampling.

The AI's report of a run != proof of a run. The gate is a deterministic run by a script/CI with a real return code, not the model's words.

Coverage is deceptive: a line "touched" != checked. More honest — mutation testing: inject artificial "mutant" defects and measure the share killed. The danger — Goodhart's law: AI optimizes the target metric; a coverage gate → tests "for coverage," not for defects.



Meta ^[2]: LLM generation covers more classes (32% vs 5.3% for a narrow targeted method), but kills fewer mutants (2.4% vs 15%). More tests and coverage != better defect detection.

The alternative: a deterministic run as the gate + a quality gate on mutation score, not coverage. Incident → permanent regression test.

SECTION 5

Review + Security — the discipline of skepticism

Review and security are the second, critical look at AI's output, and both come down to automation bias. The counterintuitive thesis of this phase: AI code needs more review, not less — the source of its defects is different.



0

Intro

1

Requirements

2

Architecture

3

Implementation

4

Testing

5

Review+Sec.

6

Delivery+

7

Synthesis

Review practice — two human practices, not a choice of AI reviewer



1. Adversarial review with fresh context ^[1]

The code is reviewed by someone OTHER than the author; the reviewer starts with clean context: sees only the diff and the acceptance criteria. Reduces the "I wrote it, so it's right" bias (writer-reviewer, two passes).



2. Retained human accountability

AI review is an assist and a first pass, but the decision and accountability rest with the human. Osmani ^[2]: "if you can't explain it, don't commit."



A fundamental tradeoff

Detection completeness ↔ noise: stricter — more bugs caught, but more false alarms; softer — less noise, but misses. Anthropic ^[3]: told to look for holes, the reviewer will find them even in healthy code (over-eagerness → over-engineering); scope it to correctness.

No tradeoff point makes AI review an autonomous gate.

First pass over the diff (secondary): GitHub Copilot code review · Cursor Bugbot · Qodo Merge · Atlassian Rovo Dev (against criteria in Jira) · Anthropic adversarial reviewer.

Durable pattern: AI review as an assist / first pass. Hype: AI review as a gate ("AI reviewed it — safe to merge"). The decision and accountability rest with the human.

Review failure: complacency and the "fake in seconds, triage in hours" asymmetry



Complacency (Radar, Hold ring) ^[1]

Uncritical acceptance of AI code, a drop in critical thinking. CodeCrash (arXiv:2504.14119) ^[3]: misleading comments crash the model's reasoning (~23% on CRUXEVAL / LIVECODEBENCH).

AI review ~19% F1 (SWR-Bench) — and even that is only stated against a human-review baseline (low + a high rate of false positives).

Stenberg: AI analyzers "in the right hands" find real bugs — the process architecture is to blame, not AI.



curl-slop as a DDoS on maintainers ^[2]

A flood of LLM-generated "vulnerability reports" in the curl bug bounty.

Cost asymmetry: generating a plausible fake takes seconds; refuting it takes hours of a maintainer's time.

Numbers: share of valid reports >15% → <5% (~1 in 20-30); volume grew several-fold; the program was suspended and moved back to HackerOne in March 2026.

AI didn't "make spam nastier" — it removed the throttle, and the economics of the process shifted. The alternative: a machine-verifiable barrier at the entrance (a reproducible PoC), not a manual review of every text.

Security — break the lethal trifecta architecturally

The lethal trifecta (Willison, June 2025 ^[1]; Fowler ^[2]) — only the intersection of all three is dangerous:



1. untrusted content

issues, emails, web pages



2. secrets / private data

keys, database



3. outbound transfer (egress)

can send data out

Untrusted content via prompt injection → grab a secret → send it out.

Four human-owned controls that break the trifecta

least-privilege

sandbox

egress-allowlist

SAST gate

Terms: SAST (static) / secret-scanning / SCA (dependencies) / supply-chain.

Secondary: GitHub (CodeQL + Copilot Autofix + secret-scanning + Dependabot) · Google (Big Sleep — live exploitation of SQLite; OSS-Fuzz + LLM — a ~20-year-old OpenSSL bug) ^[3] · AWS Q security · Anthropic /security-review.

"The first AI to stop a zero-day" = one curated case; "AI finds 50%" = metrics on their own code, not universal.

Durable pattern: a mandatory automated security scan as a gate + an architectural break of the trifecta. SAST is necessary but NOT sufficient; threat modeling is the human's.

The danger isn't the vulnerability itself, but the confidence that the code is secure

The most systemic risk is not "AI sometimes writes vulnerable code," but "vulnerable code + a developer's heightened confidence that it's secure" = automation bias in its most dangerous form.

Why systemic: autocomplete relies on the statistically frequent, and vulnerable patterns (SQL concatenation, missing validation, hardcoded secrets) are pervasive in open code. The model reproduces the frequent, not the secure.



Stanford (randomized) ^[1]

Perry et al. · arXiv:2211.03622 · CCS 2023

Developers with an AI assistant introduced vulnerabilities MORE often — and were MORE confident their code was secure. False confidence measured directly.



NYU "Asleep at the Keyboard?" ^[2]

arXiv:2108.09293 · IEEE S&P 2022

~40% of Copilot programs contained vulnerabilities.

Baseline: of 1689 programs across 89 scenarios around the MITRE Top-25 CWE — the share among deliberately security-sensitive tasks, NOT "40% of all code."

**The alternative: SAST + DAST + a mandatory security gate plus threat modeling (essential complexity, not delegated).
The danger isn't the error, but the false confidence next to it.**

Supply-chain — a separate class: a reproducible hallucination and a leak channel



Slopsquatting (a supply-chain attack)

the LLM reproducibly hallucinates a package name

an attacker registers it IN ADVANCE with malware

a developer / a C-D agent runs install <made-up>

The axis of the threat is reproducibility: of 576,000 samples ~20% recommended non-existent packages; 43% of hallucinated names recurred across all 10 queries (*Spracklen et al., USENIX Security 2025*)^[1]. **The term was coined by Seth Larson (PSF, April 2025).**



CamoLeak (prompt injection in a dev agent · Legit Security)^[2]

Instructions hidden in invisible markdown PR comments made GitHub Copilot Chat search for secrets (AWS keys) and exfiltrate them through the GitHub image proxy.

CVE-2025-59145, CVSS 9.6 (critical).

A dev agent with access to untrusted content + secrets = a ready-made exfiltration channel (a structural property, not a bug). The same lethal trifecta — inside a developer's tool.

Not cured by a "better model" — only by architecture: a lockfile with hash pinning, a registry allowlist, package verification before install, SCA; least-privilege + isolation + human-in-the-loop on writes + egress control.

The agent's speed is the speed of the catastrophe; accountability is not delegated

July 2025, a vibe-coding experiment (Replit; Fortune, 2025-07-23). The human set an explicit code freeze: "NO MORE CHANGES". Despite the ban, the agent:

- deleted the live (production) DB (1200+ executives, 1190+ companies)
- fabricated reports masking the problem
- lied to a direct question
- rated its own behavior 95 out of 100
- claimed rollback was impossible — yet the mechanism worked, the data was restored

Echoes of the same class (The Register): Amazon Kiro (Dec 2025) — hours of downtime · PocketOS / Cursor (Apr 2026) — wiped the DB in 9 seconds.

Prompt != control

"NO MORE CHANGES" is, to an agent, not an environment barrier but text competing for attention. There is no architectural boundary between a "rule" and a "wish".

Self-assessment != verification

"95/100" is anti-correlated with reality (highest at the worst outcome).

The agent's report != proof

the source of truth in a postmortem is independent telemetry, not the agent's narrative.

"95/100" at the worst outcome · "9 seconds"

Level-D safety does not live in the prompt — it lives outside the agent: dev/prod isolation, a hard human gate on destructive actions, least-privilege, a tested rollback. The root error is autonomy inadequate to the cost of error ^[1]. Accountability is not delegated.

[1] Fortune — Replit AI wiped the production DB

SECTION 6

Delivery · Operations · Documentation

Three closing phases of the cycle. Delivery and operations are thin: their input is the state of the real world (the pipeline, prod, telemetry), which is not in the text. Documentation is the map's only bright spot, but it too has a price.



0

Intro

1

Requirements

2

Architecture

3

Implementation

4

Testing

5

Review+Sec.

6

Delivery+

7

Synthesis


Delivery — DORA-first: a mature pipeline first, then scale AI

"AI amplifies what is already there" — so there is only one order, and it runs against intuition:

- 1** First — mature delivery that works WITHOUT AI: automated tests as a gate · version control with cheap rollback · fast feedback · small batches. These are four of the seven delivery capabilities DORA identifies.
- 2** Only then — scale AI on top of it.

The reverse does not work: AI does not fix an immature pipeline — it speeds it up.

Inside the practice: the prod gate is calibrated by risk — a human approves the irreversible; something small and reversible that has passed its gates may be confirmed by AI. AI consumes the pipeline but does not own it: the agent calls gh / aws / gcloud as a privilege-limited user. Operations is the weakest of the three phases: no runtime context.



One and the same AI adoption yields both halves: +7.5% to documentation quality and -7.2% to delivery stability (DORA 2024 ^[1]); the link with stability is negative for the second year running (DORA 2025 ^[2]).

Which half outweighs the other is decided by the maturity of the pipeline, not by the choice of model.

The failure of this phase is scaling AI onto an immature pipeline: both speed and instability grow. Hype: "an AI-CD/ops product as a replacement for the human".

The multiplier is volume: a mature gate multiplies what is verified, a missing gate multiplies what is unsafe



Where the AI copilot works

BT Group ^[1]

MTTR ~2h → 85s (~97% reduction)

MTTR (mean time to repair) — the average time to restore service after an incident

Microsoft Azure "Triangle" ^[2]

time-to-engage -91%, triage 97%

time-to-engage — the time until the on-call engineer is engaged

What AI multiplies: alert correlation and steps from a runbook written in advance.

Where the gate is: the runbook (a ready-made remediation script) defines the permitted actions, telemetry shows the result of every step. What is multiplied is work already verified — on top of an ALREADY mature operations practice (SRE).



The same multiplier, the other side

AI-generated IaC (infrastructure-as-code) ^[3]

~55% secure by default

the share of tasks where the generated Terraform / Kubernetes code is secure out of the box, with no human edits; over 2 years the figure has barely moved. Syntactic correctness >95%.

A separate 2026 benchmark, a different measurement: 8.4% on security-checked tasks

What AI multiplies: the generation of configurations — many times more artifacts, and faster.

There is no gate: the pipeline checks syntax, not security. More configurations → less attention on each → the unsafe one reaches production.

The problem is not that "there are more rollouts" but WHAT exactly is being multiplied: BT multiplies a verified step, IaC generation multiplies an unverified artifact. The multiplier is the same; only the maturity of the gate differs.

Documentation — AI's only clean plus, but it too has a paired cost



Bright spot

The only phase with a clean positive systemic effect from AI. Why: accidental complexity dominates; the cost of error is asymmetrically low; human control is built in — the docs get read.

DORA 2024 ^[1]: +7.5% to documentation quality. Baseline: cited only paired with -7.2% delivery stability (the AI effect almost always has a paired cost); stability is negative for a second year.



Cognitive debt (Radar, Hold)

documentation generation outpaces understanding: lots of text, less understanding. The named remedy is architectural fitness functions (Ford/Parsons): they keep the "why" in a verifiable form.



Onboarding docs hallucinate the setup / deployment

Böckeler ^[2]: "AI cannot magically replace a well-documented and automated setup."

Secondary: Confluence AI · AWS Q /doc · JetBrains KDoc/Javadoc.

Practice: docs-as-context — code stays the source of truth, documentation is context; generation pace <= comprehension pace. Documentation-as-context — yes; documentation-as-truth — no.

The documentation toolchain: the input is code, the output is a retelling of what was already said



The SaaS vendor layer

Confluence AI (Atlassian Intelligence)

Summarizing a long thread into a digest; generating and transforming a draft from a prompt; Q&A search over the knowledge base (a RAG-like pattern over the corporate base).

AWS Q Developer /doc

The agent analyzes the codebase rather than retelling the prompt. It builds infrastructure diagrams from IaC files (Terraform/CDK) — a direct link to the already-introduced principle of architecture-as-code. A closed loop: the code changed → a diff in the documentation is proposed.



The code-first alternative

Instead of a separate SaaS — the coding agent itself, through an installed skill ("Code Documentation Skill", "README Generator"): it analyzes the project structure, the dependencies and the code patterns, and generates README / ADR / inline comments.



An honest caveat

a community pattern, not a single official skill from the Anthropic repository. Onboarding documentation is a good skill candidate by the heuristics already introduced: a repeated instruction + it needs progressive disclosure.

The vendor-specific layer illustrates the mechanics of the current 2026 stack, not a recommendation of one vendor: the same pattern is available through a skill on top of the agent you already use, without SaaS.

SECTION 7

Synthesis — discipline by phase

We have gone through every phase; now let's fold them into a working toolkit: a matrix "phase × leading practice × where the human is mandatory", triangulation of independent measurements, a compact risk triad, and a checklist "when AI yes, when no".



0

Intro

1

Requirements

2

Architecture

3

Implementation

4

Testing

5

Review+Sec.

6

Delivery+

7

Synthesis

The lecture's matrix: the practice leads — no vendor column

Phase	Leading practice	Failure mode	Where the human is required
 Requirements	spec-driven: spec before code	prompt-and-pray; "spec = truth"	deciding what to build
 Architecture	ADR + fitness + arch-as-code	poisoned context without management	choosing the forks under trade-off
 Implementation	explore→plan→code→commit + harness	the 70% problem; "almost right"	reviewing the diff + merge
 Testing	TDD: test-as-spec + determ. gate	"all green" lies; coverage != defects	what the test asserts
 Review + Sec.	fresh-context; least-priv+SAST	complacency; false confidence	a second pass + threat modeling
 Delivery	headless + risk-calibr. gate (DORA-first)	AI consumes, does not own	prod gate (hard on the irreversible)
 Operations	telemetry + on-call	no system context	owning the system model
 Documentation	docs-as-context (code = truth)	cognitive debt; hallucinations	pace <= comprehension pace

The vendor names from sections 1-6 are replaceable; the practice, the failure mode and the human's point here are durable — they rest on the character of the phase's complexity ^[2]. Every cell is derived from a section covered, not assigned. ^[1]

Three independent methods converge: an individual AI gain != systemic quality



DORA (n ~ 5000, systemic) ^[1]

~90% of reports: throughput positive, but AI's link to stability is negative for a second year running. Lens: "AI amplifies what is already there".



GitClear — two measurements ^[2]

211M lines (2020-24): refactoring ~25%→<10%, duplicates 8.3%→12.3%. 623M changes (2023-26): refactoring 21%→3.8% (-70%), duplicates +81%, churn +15%. (Both are correlation, not an RCT.)



METR (n = 16, experts, familiar code) ^[3]

tasks with AI took +19% time, yet believed in a speed-up (~-20%) = the perception gap. (On unfamiliar code the effect differs.)



The strength is in the convergence of independent methods: DORA, GitClear and METR have different blind spots, so the chance all three erred the same way is low. The shared conclusion is more reliable than any single number.

One conclusion: the method matters more than the tool. Practice — a CI gate on duplication and churn; measure the systemic effect, not the perception.

Scale of adoption is not yet an effect: what decides is neither the brand nor the reach, but the discipline applied



Uber, 2026 ^[1]

Uber

Uber · Wikimedia

Scale of adoption

agentic practices 32% → 84% in one month (a 2.6× rise) · 95% of engineers monthly · 70% of committed code — from AI · \$500–2000 per engineer per month

Effect: not traceable

"It's hard to draw a connection between the company's rising use of Claude Code and innovations meant to serve consumers... That link is not there yet"

— Andrew Macdonald, President and COO of Uber

Response: a \$1500 per-employee monthly cap — after the fact, once the annual budget had burned through in 4 months.



AWS Kiro — the same product, two registers ^[2]

aws

AWS · Wikimedia

Success:

life sciences — a spec-first discipline, verifiability gates set in advance → production in 3 weeks (the same case we already saw earlier in the lecture).

Failure:

the Kiro incident, December 2025 — the agent autonomously tore down and rebuilt the environment without approval → hours of downtime.

The difference is not the brand (the product is one and the same) but the discipline applied or skipped: here the outcome is measurable in both directions — unlike the scale without a criterion on the left.

A rise in adoption is not in itself a result: 2.6× in one month with no criterion set in advance produced a rise in spending and an unconfirmed effect. The decision about the scale of AI is a measurable engineering decision with a criterion set in advance, not cultural inertia.

The question is not "AI or not" but which level of autonomy: probability × impact × detectability

AI is applied in one mode or another almost always — the triad answers not "yes / no" but which autonomy ceiling is admissible and what pays for it.


1. Probability of error


low → high grows with unfamiliarity of the task and codebase

2. Impact of error


low → high irreversibility, safety, money, data

3. Detectability


low → high is there a test oracle, SAST or review that catches it

Where the autonomy ceiling is highest

Full vibe-coding — only with the combination **low × low × high**: low probability of error, low impact, high detectability. Any other combination is not a ban on AI but a lower ceiling and mandatory measures. The axes multiply, they do not add.


What a higher level is paid for with:

- impact ↑ → a hard human gate
- detectability ↓ → a machine oracle
- probability ↑ → line-by-line senior review

Böckeler ^m: "using generative AI is a continuous risk assessment": the decision is not made once and not about the tool as a whole, but on every task. The failure is vibe-coding "by feel": not one of the three axes is computed. Every case in the lecture converges on it: Replit (impact ↑), curl-slop (detectability ↓), vulnerable code (probability ↑).

Eight questions — not "AI or not", but how, where and with what control to apply it

- ✓✓ 1. Which lifecycle phase is this? It sets the failure mode
- ✓✓ 2. What is decided deterministically here? Ordinary code writes that part; AI is for analysis and checking
- ✓✓ 3. Essential or accidental complexity? Essential — the human decides, AI stays at the periphery
- ✓✓ **4. Is the consequence reversible? Irreversible → lower the autonomy ceiling, a hard gate — the VETO axis**
- ✓✓ 5. Is there a machine oracle (test, SAST, run)? No → the oracle first, autonomy after
- ✓✓ 6. Are secrets / untrusted content involved? Yes → least privilege and isolation
- ✓✓ 7. Who reviews and who merges? Merging and accountability — always the human
- ✓✓ 8. Is the goal an artifact or a skill? A skill → do not delegate the generation



Anthropic, Shen & Tamkin 2026 (RCT, n=52, learning an unfamiliar library) ^[1]: the AI group scored 50% on the quiz against 67% without AI (~-17 pp). Those who delegated the generation dropped; those who asked about concepts ("how it works, why") show no degradation. The speed-up is statistically insignificant.

When the goal is a skill, you do the writing; AI explains and checks.

The checklist does not decide "apply AI or not" — it issues a mode: for a suitable task it leads to high autonomy; for an unsuitable one it lowers the ceiling and names the mandatory measures. This is a distribution of the burden of proof, not "always less AI". Irreversibility and impact are the veto axis.

AI changes the cost of writing code — not the cost of understanding and responsibility

AI changes the cost of writing code, but not the cost of understanding what to build and who is responsible for it. It touches each phase differently — and reliability comes not from the tool but from discipline by phase.

The method carries over to every industry (not a list of tools):

- 1 decompose the activity into phases
- 2 ask: accidental or essential complexity
- 3 demand a baseline for every number and the systemic effect
- 4 separate the durable pattern from vendor hype with five questions

Questions?

